

Quiz 6: Similar Items

Name: GRADER ID: _____

Input: 10 million documents in English in the database and 1 new document

Output: in the database, all the documents that are similar to the new document with 80% Jaccard similarity and their Jaccard similarity to the new document

Describe your process to generate the output from the input using k-shingles, minhash, and LSH. You need to describe your input, process, and output using examples.

1. How to generate features describing each document using 10-shingles (1 pt)

For each document generate unique 10-shingles sets and then converts them into integers as features.

	D ₁	D ₂	...
S ₁	0	1	0
S ₂	1	0	0
S ₃	0	0	1

2. How to efficiently generate 100 minhash signatures for each document from their 10-shingle features (2 pts)

[1 point] → Design 100 hash functions and find its 100 min hash signatures by scanning the table and indicate the row number where first '1' appears.

	D ₁	D ₂	D ₃	...	D ₁₀₀₀₀	h ₀	h ₁	h ₂	...	h ₉₉
S ₁	0	1	0	...	0	1
S ₂	1	0	1	0
S ₃	0	0	0	5

[1 point] → minhash signature

3. How to use LSH to speed up the process of comparing signatures; what are the parameters in your LSH process? (3 pts); what are the false-positives and false negatives? (1 pt) how to set the parameters to control false-positives and false negatives? (3 pts)

[1.5 POINTS] To speed up the process, ~~one~~ sample the signatures with several bands and just compare the signatures in the same bands. if similar then put into candidate pairs

[1.5 POINTS] Parameters: r: rows of each band, b: number of bands.
s: Jaccard similarity, t: threshold.

[0.5] False positives: Items are dissimilar but hashed to same band/bucket.

[0.5] False negatives: Items are similar but hashed to different band/bucket.

[1.5 POINTS] r: ~~increase~~ False increase r it will increase false negative but reduces false +ve

[1.5 POINTS] b: increase b it will increase false positive but reduces false-negatives